

Exam INFO-H-417 Database systems architecture

January 2022

Mahmoud SAKR

Name:

ULB Student ID:

Q1 (25)				Q2 (25)			Q3 (15)		Q4 (25)	

Important notes

- Please make sure to sign your name with the monitor twice: when you enter the exam room, and before you leave it. Without these two signatures, you will be considered absent.
- Please make sure to write your name and student ID on this page. On every other paper, including the draft, make sure to write your student ID. A paper without your ID number is considered a cheat source. Your ID number on every paper also guarantees identifying your answers in case some staples break.
- You are allotted a maximum of **3 hours** to complete this exam.
- The exam is closed book. You are hence not allowed to use any notes/books/PC/etc. during the exam.
- Draft papers are provided by the monitor upon request. **Draft paper will not be corrected!**
- You should answer each question in the foreseen space after the question. Should this space prove to be insufficient you are allowed to use the back of the page as well.
- Mobile phones cannot be used during the whole exam duration.

STUDENT ID: _____

(b) Write an RA expression to express this query:

List all department id and name of the students who attend the Database course
Make sure not to show duplicate results

(c) Normalize/flatten the following query and give its equivalent RA expression

```
SELECT course.name
FROM course
WHERE course.id in (
    SELECT enrollment.course_id
    FROM student, enrollment
    WHERE student.id = enrollment.student_id AND
    student.name LIKE '%son');
```

STUDENT ID: _____

2. Attribute statistics for the optimizer

3. The optimization rules

4. User tables, such as course, student, enrollment in question 1

5. The SQL user functions

6. The B+tree indexes

7. User defined type

8. The database schema

(c) State whether each of the following points are true about images in System R. Give one line discussion on every point:

1. One may define multiple images on the same relation

2. One may define multiple images on the same attribute

3. Image may be declared as unique

STUDENT ID:

4. Defining a clustered image may lead to changing the ordering of the stored tuples in the physical medium

5. One may define multiple clustered images on the same relation

6. One may define multiple unique images on the same relation

7. If available, it is always more efficient to access a relation via an image

Question 3. (15 pts)

(a) Recall that in a B+tree:

- one tree node is stored in one disk block. In this exercise consider that the block size is 4096 B
- one tree node stores p pointers and $p-1$ keys

Consider a B+tree index defined over an int16 attribute. Consider also that a tree/disk pointer is 8 bytes. If the relation that contains the indexed attribute has four million tuples, how many disk accesses would be required to answer a query $key = const$, where $const$ is a literal value?

STUDENT ID: _____

(b) Discuss the truth of each of the following statements.

1. The leaf level of a B+tree index is always a dense index

2. The root level of a B+tree index is always a dense index

3. Inserting duplicate keys can result in a non-balanced B+tree

4. A B+tree index may require a bigger storage size than its relation

5. A B+tree index on a string/text attribute cannot be used to answer queries in the form of key *LIKE* pattern

Question 4. (25 pts)

(a) Give a short answer on the following.

1. What is an operator class in PostgreSQL? Give its importance.
2. How is it possible for PostgreSQL to process user types, without knowing them in advance?
3. How is a materialized view different from a table? How does a materialized view affect query performance.
4. How does tuple insertion happen over a distributed table using hash-distribution?
5. What is the importance of replicated tables in distributed databases? Why not distribute all tables?

STUDENT ID: _____

STUDENT ID: _____