



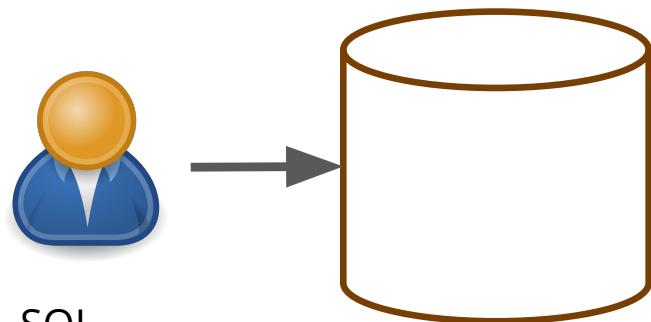
INFOH417 Database System Architectures

Mahmoud SAKR <mahmoud.sakr@ulb.be>

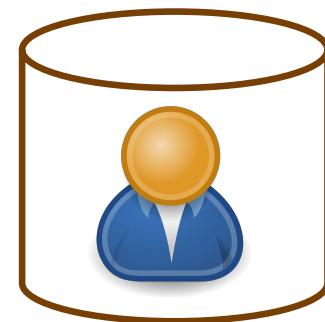
École polytechnique de Bruxelles

2022

What is this course for ?



SQL
Python
Web
...



Storage
Access control
Optimization
Distribution
...

Course Goals

- Understanding the query optimization and execution cycle
- Improving slow queries
- Describing the common index structures, knowing their capabilities and shortcomings
- Understanding cost based optimization, and the associated statistics and estimation methods
- Describing and being able to implement Abstract Data Types in extensible database systems
- Describing data and query distribution mechanisms, and being able to configure and run a distributed database system

Course Topics

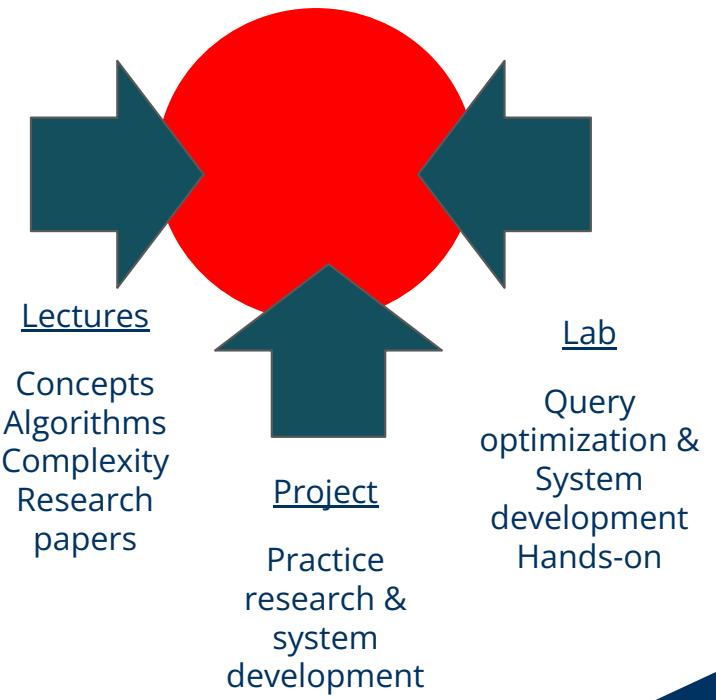
- Query execution
- Refreshing SQL and relational Algebra
- Cost-based query optimization
- Indexes
- Crash recovery
- Concurrency control
- Distributed databases

Prerequisites

- Relational databases
- SQL
- Relational Algebra
- General programming skills

Course Organization

- Lecture schedule and room on timedit: Check regularly for schedule updates
- Lectures by mahmoud.sakr@ulb.be
- Teams broadcasting and Recording will be there without any guarantees. Don't depend on it, many tech problems can happen.
- Lab sessions by: Mariana Duarte<mariana.machado.garcez.duarte@ulb.be>
 - Install PostgreSQL, any version, use same as the advanced DB course
 - Only 6 lab sessions, see UV for the schedule
- Grading
 - 2x Group projects, 4 members, 40% in total
 - Written exam, 60%
- Course notes, please enroll in [Université virtuelle](#)



Recommended Readings

- A mixture of book chapters and research papers, which will be identified per lecture

Query Planning: Translating SQL into Relational Algebra

The set of attributes
of a relation is a
schema

attributes
with a certain datatype

A	B	C	D
1	2	3	4
1	2	3	5
3	4	5	6
5	6	3	4

tuples

Refreshing the Relational Algebra

- Relations are tables whose columns have names, called attributes
- The set of all attributes of a relation is called the schema of the relation
- The rows in a relation are called tuples
- A relation is set-based if it does not contain duplicate tuples.
- It is called bag-based otherwise.
- A Relational Algebra (RA) operator takes as input 1 or more relations

and produces as output a new relation

process the
data from the
relations

playing with relations
to create a new relation

) allows for duplicate
tuples
defends
on the
implementation
 \Rightarrow defends on
the database
server
we are using

the whole table
is a relation

set semantic will remove redundancies
 bag semantic will not

select-from-where

SQL:

```

SELECT movieTitle
FROM StarsIn S, MovieStar M
WHERE S.starName = M.name AND M.birthdate = 1960
  
```

RA?

joining S and M

attributes
to output

relations to
take input data from

conditions

$\pi_{\text{movieTitle}} \left(\sigma_{n.\text{birthdate} = 1960} \left(\chi_{\Pi}(\text{MovieStar}) \right) \times_{M.\text{name} = S.\text{starName}} \chi_S(\text{StarsIn}) \right)$

project from & attributes → to only movieTitle

select-from-where

SQL is a declarative language,
you describe what you want not how to achieve it
not useable by computer

SQL:

```
SELECT movieTitle  
FROM StarsIn S, MovieStar M  
WHERE S.starName = M.name AND M.birthdate = 1960
```

↳ RA is a procedure

you describe how to get the result

) can be executed by a computer

RA ?
projection

select
 $\prod_{\sigma_{S.starName=M.name \text{ and } M.birthdate=1960}} \pi_{\text{movieTitle}} (\rho_M(\text{MovieStar}))$

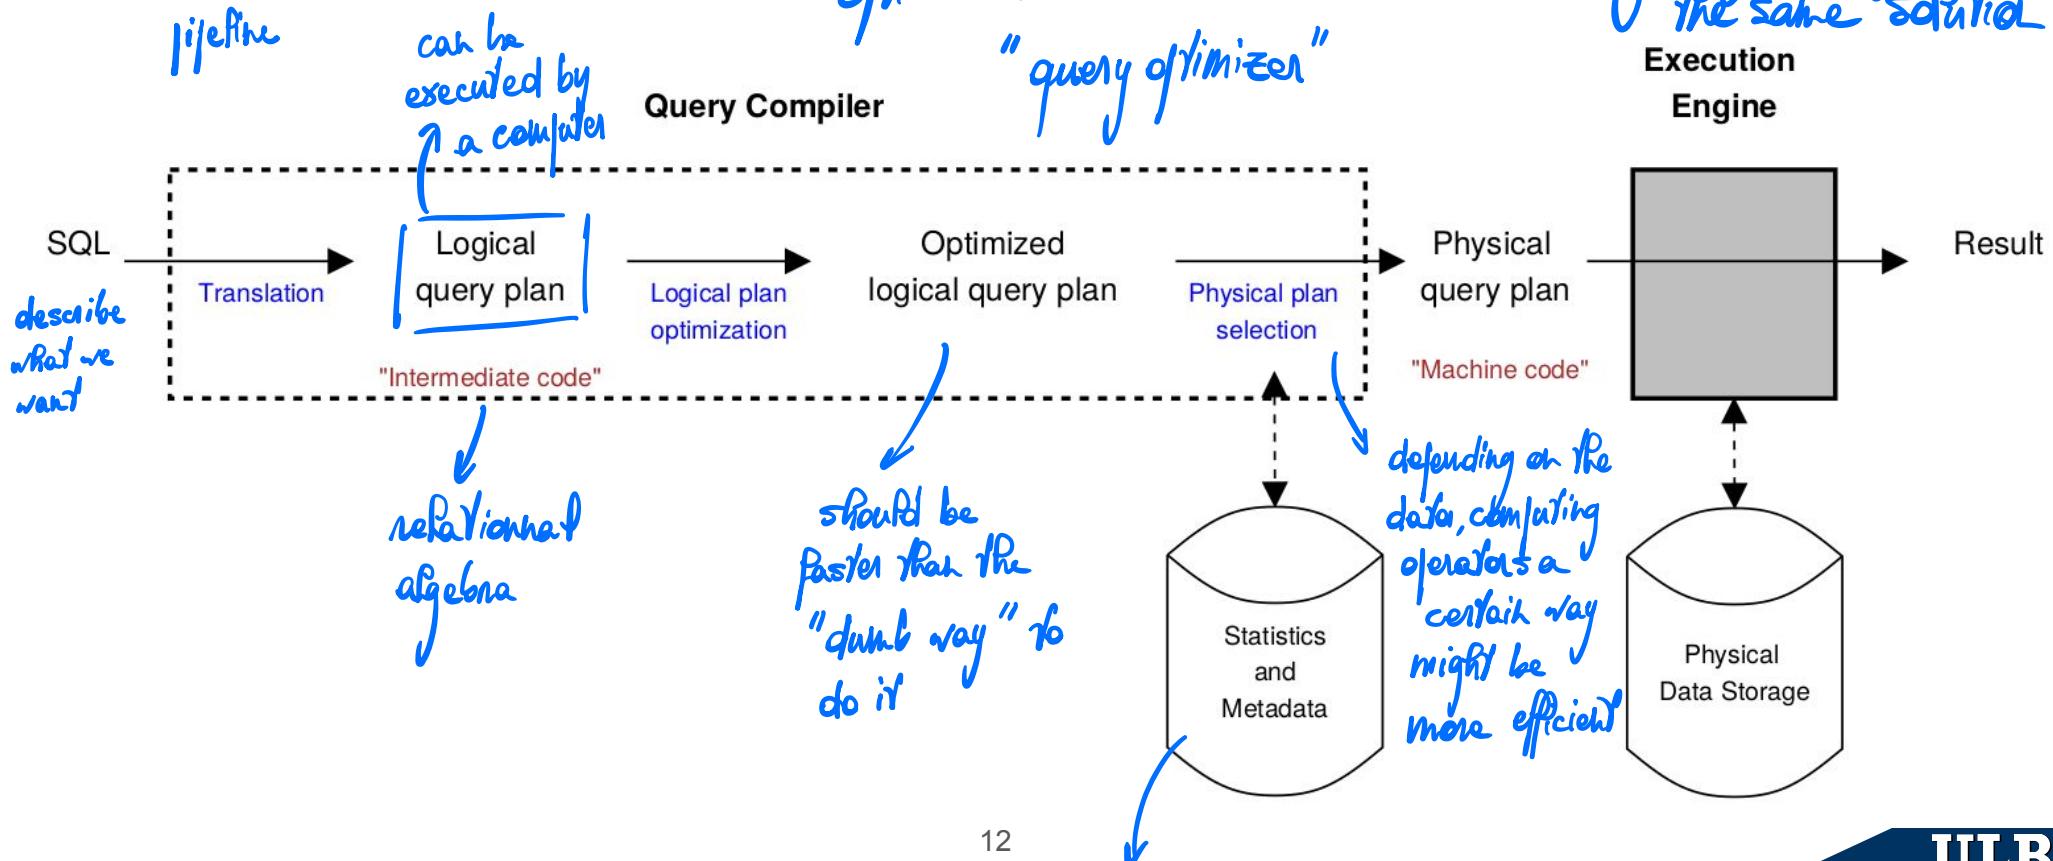
Other translations ?

→ rename

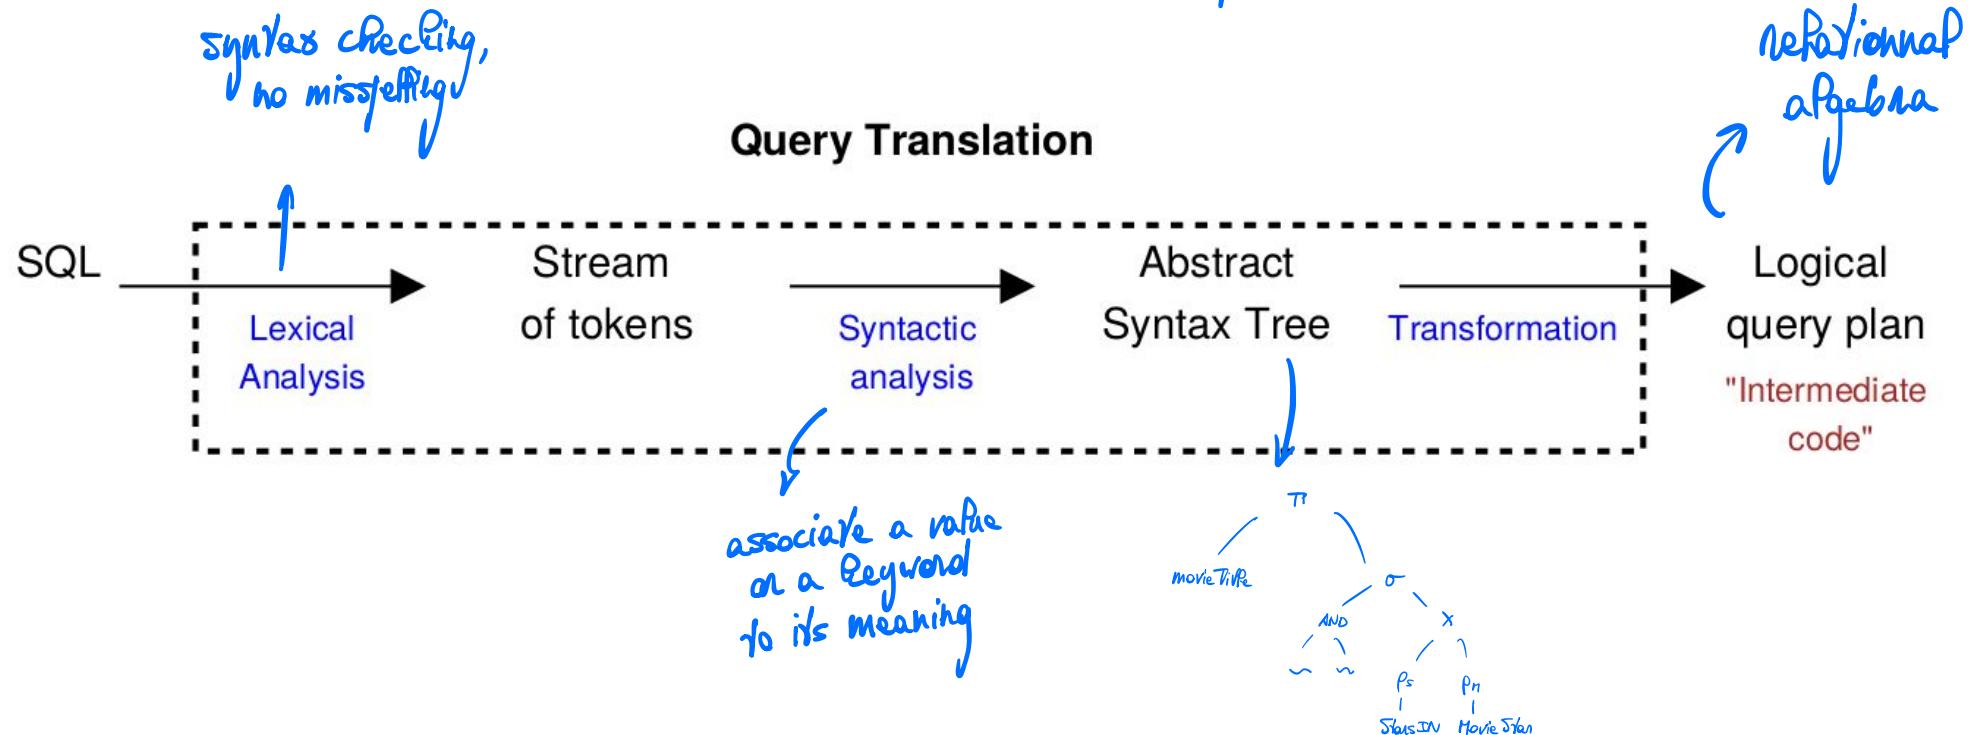
$(\rho_S(\text{StarsIn}))$
rename

translation
from SQL to RA
is required

Query Planning



Query Translation



The Extended Relational Algebra

Operator		Operator	
\cup	Union	\bowtie	Natural join
\cap	Intersection	$\bowtie_{B=C}$	Theta join
-	Difference	$\bowtie_{B=C}$	Left outer join
$\sigma_{A>=3}$	Selection	$\bowtie_{B=C}$	Right outer join
$\pi_{A,C}$	Projection	$\bowtie_{B=C}$	Full outer join
\times	Cartesian product	γ	Aggregation
ρ_T	Rename	$A, \min(B) \rightarrow D$	Assignment

Union \cup / Intersection \cap / Difference -

A	B
1	2
3	4
5	6

 \cup

A	B
3	4
1	5

 $=$

A	B
1	2
3	4
5	6
1	5

Set-based \rightarrow *remove duplicates*

Bag-based

A	B
1	2
3	4
5	6
1	5
3	4

- Input relations must have the same **schema** (same set of attributes)
- Historically speaking, relations are defined to be sets of tuples: duplicate tuples cannot occur in a relation.
- In practical systems, however, it is more efficient to allow duplicates to occur in relations, and only remove duplicates when requested. In this case relations are bags.

Selection

$\sigma_{A>=3}$

A	B
1	2
3	4
5	6

=

A	B
3	4
5	6

select tuples
where $A >= 3$

Projection

$\Pi_{A,C}$

A	B	C	D
1	2	3	5
3	4	3	6
5	6	5	9
1	6	3	5

Set-based

A	C
1	3
3	3
5	5

can create redundancies

Cartesian Product

A	B
1	2
3	4

\times

C	D
2	6
3	7
4	9

=

A	B	C	D
1	2	2	6
1	2	3	7
1	2	4	9
3	4	2	6
3	4	3	7
3	4	4	9

The two input relations
can be of different
schema

generate all
the possibilities

but concatenation
column-wise
 \Rightarrow must have two
attributes with
the same name

it does a row-wise
cross product of every
tuple

Input relations must have disjoint schema (disjoint set of attributes), otherwise rename first

Natural Join

equijoin on attributes
with the same name

"common attributes"

\bowtie

A	B	
1	2	
3	4	

\bowtie

B	D	
2	6	
3	7	
4	9	

=

A	B	D
1	2	6
3	4	9

↳ if multiple attributes with the
same name it will join
by all of them

Natural Join

A	B	C	D
1	2	2	6
3	4	3	7

⊗ =

A	B	C	D
1	2	2	6
1	2	3	7
1	2	4	9
3	4	2	6
3	4	3	7
3	4	4	9

relations with no
shared attributes

Same as cartesian product

↳ equijoin of everything on everything
→ same as cross-product

Theta Join

A	B
1	2
3	4

 $\bowtie_{B=C}$

natural join
variant in the case
of no common attributes,
we can still do an equijoin
on a user specified condition

C	D
2	6
3	7
4	9

=

A	B	C	D
1	2	2	6
3	4	4	9

gives flexibility

but can be done using
other operators

→ not an
essential operator

renaming
+ natural join

Renaming

$$\rho_T \quad \begin{array}{|c|c|} \hline A & B \\ \hline 1 & 2 \\ \hline 3 & 4 \\ \hline 5 & 6 \\ \hline \end{array} \quad = \quad \begin{array}{|c|c|} \hline T.A & T.B \\ \hline 1 & 2 \\ \hline 3 & 4 \\ \hline 5 & 6 \\ \hline \end{array}$$

Renaming specifies that the input relation (and its attributes) should be given a new name.

Relational Algebra Expressions

Built using relation variable, AND

RA operators

$$\sigma_{\text{length} >= 100}(\text{Movie}) \bowtie_{\text{title} = \text{movietitle}} \text{StarsIn}$$

Write the equivalent SQL

SELECT * FROM Movie, StarsIn
WHERE length >= 100
AND title = movieTitle

SELECT * FROM Movie
JOIN StarsIn ON title = movieTitle
WHERE length >= 100

The Extended Relational Algebra

↳ but doesn't make a big difference
since we are still in the SQL
world

Add more operators

Extended projection

allows renaming

Π
 $A, C \rightarrow D$
combined
select and
renaming

A	B	C	D
1	2	3	5
3	4	3	6
5	6	5	9
1	6	3	5

= Set-based

A	D
1	3
3	3
5	5

The Extended Relational Algebra

Add more operators

Grouping

$\gamma_{A, \min(B) \rightarrow D}$
group by A

A diagram illustrating the grouping operation. On the left, a 6x3 table is shown with columns labeled A, B, and C. The rows contain the following data:

	A	B	C
1	1	2	a
1	1	3	b
2	2	3	c
2	2	4	a
2	2	5	a

On the right, the result of the grouping operation is shown as a 2x2 table with columns labeled A and D, containing the values 1 and 2 respectively.

=

	A	D
1	1	2
2	2	3

Translating SQL into Relational Algebra

In the examples that follow, we will use the following database:

- Movie(title: string, year: int, length: int, genre: string, studioName: string, producerC#: int)
- MovieStar(name: string, address: string, gender: char, birthdate: date)
- StarsIn(movieTitle: string, movieYear: string, starName: string)
- MovieExec(name: string, address: string, CERT#: int, netWorth: int)
- Studio(name: string, address: string, presC#: int)

select-from-where-groupby

SQL:

```
SELECT movieTitle, count(S.starName) AS numStars  
FROM StarsIn S, MovieStar M  
WHERE S.starName = M.name  
GROUP BY movieTitle
```

RA ?

$\delta_{movieTitle, count(\dots) \rightarrow numStars} (MovieStar \bowtie_{starName = Name} StarsIn)$

select-from-where-groupby

SQL:

```
SELECT movieTitle, count(S.starName) AS numStars  
FROM StarsIn S, MovieStar M  
WHERE S.starName = M.name  
GROUP BY movieTitle
```

RA ?

$\gamma_{M.movieTitle, \text{count}(S.starName) \rightarrow \text{numStars}} ($

$\rho_S(\text{StarsIn}) \bowtie_{S.starName=M.name} \rho_M(\text{MovieStar}))$

select-from-where-groupby-having

```
SELECT movieTitle, count(S.starName) AS numStars  
FROM StarsIn S, MovieStar M  
WHERE S.starName = M.name  
GROUP BY movieTitle  
HAVING count(S.starName) > 5
```

RA ?

$\sigma_{\text{numStars} > 5} (\gamma_{\text{movieTitle}, \text{count}(S.\text{starName}) \rightarrow \text{numStars}} ($
 $\text{StarsIn } \pi_{\text{starName} = \text{name}} \text{ Movie}))$

select-from-where-groupby-having

```
SELECT movieTitle, count(S.starName) AS numStars  
FROM StarsIn S, MovieStar M  
WHERE S.starName = M.name  
GROUP BY movieTitle  
HAVING count(S.starName) > 5
```

RA ?

$$\sigma_{\text{numStarts} > 5} (\gamma_{M.\text{movieTitle}, \text{count}(S.\text{starName}) \rightarrow \text{numStars}} ($$

$$\rho_S(\text{StarsIn}) \bowtie_{S.\text{starName} = M.\text{name}} \rho_M(\text{MovieStar})))$$

Subqueries

A. `SELECT *
FROM huge
WHERE c1 IN
(SELECT c1 FROM tiny)`

V.S.

B. `SELECT *
FROM huge h, tiny t
WHERE h.c1=t.c1`

Which query is better?

①

project before \Rightarrow reduce
number of columns

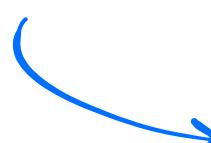
A `huge.c1 huge.c2`

B `huge.c1 huge.c2 tiny.c2`

②

if multiple c1 with the same
value in tiny relation,
B will cause more tuples
than A

same result is self semantic, but...



the query must be done for every
tuple in huge

\Rightarrow More room for optimization
in B

PostgreSQL Source Code git master

Main Page	Namespaces ▾	Data Structures ▾	Files ▾	
-----------	--------------	-------------------	---------	--

- ▶ foreign
- ▶ jit
- ▶ lib
- ▶ libpq
- ▶ main
- ▶ nodes
- ▶ optimizer
 - ▶ geqo
 - ▶ path
 - ▶ plan
 - ▶ analyzejoins.c
 - ▶ createplan.c
 - ▶ initplan.c
 - ▶ planagg.c
 - ▶ planmain.c
 - ▶ planner.c
 - ▶ setrefs.c
 - ▶ subselect.c
- ▶ prep
- ▶ util
- ▶ parser

```
645  /*
646   * If there is a WITH list, process each WITH query and either convert it
647   * to RTE_SUBQUERY RTE(s) or build an initplan SubPlan structure for it.
648   */
649  if (parse->cteList)
650    SS_process_ctes(root);
651
652  /*
653   * If the FROM clause is empty, replace it with a dummy RTE_RESULT RTE, so
654   * that we don't need so many special cases to deal with that situation.
655   */
656  replace_empty_jointree(parse);
657
658  /*
659   * Look for ANY and EXISTS SubLinks in WHERE and JOIN/ON clauses, and try
660   * to transform them into joins. Note that this step does not descend
661   * into subqueries; if we pull up any subqueries below, their SubLinks are
662   * processed just before pulling them up.
663   */
664  if (parse->hasSubLinks)
665    pull_up_sublinks(root); ← tries to flatten subqueries
666
667  /*
668   * Scan the rangetable for function RTEs, do const-simplification on them,
669   * and then inline them if possible (producing subqueries that might get
670   * pulled up next). Recursion issues here are handled in the same way as
671   * for SubLinks.
672   */
673  preprocess_function_rtes(root);
674
675  /*
676   * Check to see if any subqueries in the jointree can be merged into this
677   * query.
678   */
679  pull_up_subqueries(root); ←
680
```

Subquery processing and transformations

Subqueries are notoriously expensive to evaluate. This section describes some of the transformations that Derby makes internally to reduce the cost of evaluating them.

[Materialization](#)

Materialization means that a subquery is evaluated only once. There are several types of subqueries that can be materialized.

[Flattening a subquery into a normal join](#)

[Flattening a subquery into an EXISTS join](#)

[Flattening VALUES subqueries](#)

[DISTINCT elimination in IN, ANY, and EXISTS subqueries](#)

[IN/ANY subquery transformation](#)

Parent topic: [Internal language transformations](#)

Related concepts

[Predicate transformations](#)

[Transitive closure](#)

Subqueries

We can always normalize subqueries to use only **EXISTS** and **NOT EXISTS** [Van den Bussche, Vansummeren]^{1,2}

```
SELECT movieTitle FROM StarsIn  
WHERE starName IN (SELECT name  
                    FROM MovieStar  
                    WHERE birthdate=1960)
```

⇒

```
SELECT movieTitle FROM StarsIn  
WHERE EXISTS (SELECT name  
                  FROM MovieStar  
                  WHERE birthdate=1960 AND name=starName)
```

1 Only valid for set-based Relations

2 https://cs.ulb.ac.be/public/_media/teaching/infoh417/sql2alg_eng.pdf

Subqueries

We can always normalize subqueries to use only **EXISTS** and **NOT EXISTS** [Van den Bussche, Vansummeren]^{1,2}

```
SELECT name FROM MovieExec  
WHERE netWorth >= ALL (SELECT E.netWorth  
                         FROM MovieExec E)
```

⇒

```
SELECT name FROM MovieExec  
WHERE NOT EXISTS(SELECT E.netWorth  
                  FROM MovieExec E  
                  WHERE netWorth < E.netWorth)
```

1 Only valid for set-based Relations

2 https://cs.ulb.ac.be/public/_media/teaching/infoh417/sql2alg_eng.pdf

Subqueries

We can always normalize subqueries to use only EXISTS and NOT EXISTS [Van den Bussche, Vansummeren]^{1,2}

```
SELECT C FROM S  
WHERE C IN (SELECT SUM(B) FROM R  
             GROUP BY A)
```

⇒ ?

```
SELECT C FROM S  
WHERE EXISTS (  
    SELECT * FROM R  
    GROUP BY A  
    HAVING SUM(B) = C )
```

1 Only valid for set-based Relations

2 https://cs.ulb.ac.be/public/_media/teaching/infoh417/sql2alg_eng.pdf

Subqueries

We can always normalize subqueries to use only **EXISTS** and **NOT EXISTS** [Van den Bussche, Vansummeren]^{1,2}

```
SELECT C FROM S  
WHERE C IN (SELECT SUM(B) FROM R  
            GROUP BY A)
```

⇒

```
SELECT C FROM S  
WHERE EXISTS (SELECT SUM(B) FROM R  
              GROUP BY A  
              HAVING SUM(B) = C)
```

1 Only valid for set-based Relations

2 https://cs.ulb.ac.be/public/_media/teaching/infoh417/sql2alg_eng.pdf

Normalization

- Before translating a query we first normalize it such that all of the subqueries that occur in a WHERE condition are of the form EXISTS or NOT EXISTS.

Correlated Subqueries

A subquery can refer to attributes of relations that are introduced in an outer query.

```
SELECT movieTitle  
FROM StarsIn S  
WHERE EXISTS (SELECT name  
               FROM MovieStar  
               WHERE birthdate=1960 AND name=S.starName)
```

context of the subquery

both are correlated
through the parameters
i.e. the attributes of the
context used in the
subquery

- The “outer” relations are called the context relations of the subquery.
- The set of all attributes of all context relations of a subquery are called the parameters of the subquery.

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName  
FROM StarsIn S, Movie M  
WHERE S.movieYear >= 2000  
AND S.movieTitle = M.title  
AND EXISTS (SELECT name  
            FROM MovieStar  
            WHERE birthdate=1960 AND name= S.starName)
```

$\pi_{S.movieTitle, M.studioName}$
 $\{ \pi_{M.title, M.studioName} [\pi_{name} ((\sigma_{birthdate=1960} (MovieStar)) \bowtie_{name=S.starName} (P_n(Movie))) \bowtie_{S.movieTitle = M.title}] \}$
 $(\sigma_{S.movieYear \geq 2000} (P_S(StarsIn))) /$
 $\pi_{S.starName, S.movieTitle}$

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName  
FROM StarsIn S, Movie M  
WHERE S.movieYear >= 2000  
AND S.movieTitle = M.title  
AND EXISTS (SELECT name  
            FROM MovieStar  
            WHERE birthdate=1960 AND name= S.starName)
```

First translate the subquery

$$\Pi_{\text{name}} \ \sigma_{\text{birthdate}=1960 \wedge \text{name}=\text{S.starName}} (\text{MovieStar})$$

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName  
FROM StarsIn S, Movie M  
WHERE S.movieYear >= 2000  
AND S.movieTitle = M.title  
AND EXISTS (SELECT name  
            FROM MovieStar  
            WHERE birthdate=1960 AND name= S.starName)
```

Fix: add the context relation and parameters

$$\Pi_{\text{name}, \text{S.movieTitle}, \text{S.movieYear}, \text{S.starName}} \sigma_{\text{birthdate}=1960 \wedge \text{name}=\text{S.starName}} (\text{MovieStar} \times \rho_S (\text{StarsIn}))$$

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName  
FROM StarsIn S, Movie M  
WHERE S.movieYear >= 2000  
AND S.movieTitle = M.title  
AND EXISTS (SELECT name  
            FROM MovieStar  
            WHERE birthdate=1960 AND name= S.starName)
```

Next, translate the **FROM clause of the outer query**

$$\rho_S(\text{StarsIn}) \times \rho_M(\text{Movie})$$

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

Synchronize both expressions by means of a join.

$$\begin{aligned} \rho_S(\text{StarsIn}) \times \rho_M(\text{Movie}) \bowtie \\ (\pi_{\text{name}, S.\text{movieTitle}, S.\text{movieYear}, S.\text{starName}} \\ \sigma_{\text{birthdate}=1960 \wedge \text{name}=S.\text{starName}} (\text{MovieStar} \times \rho_S(\text{StarsIn}))) \end{aligned}$$

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName
FROM StarsIn S, Movie M
WHERE S.movieYear >= 2000
AND S.movieTitle = M.title
AND EXISTS (SELECT name
            FROM MovieStar
            WHERE birthdate=1960 AND name= S.starName)
```

Simplify

$$\rho_M(\text{Movie}) \bowtie (\pi_{S.\text{movieTitle}, S.\text{movieYear}, S.\text{starName}} \sigma_{\text{birthdate}=1960 \wedge \text{name}=S.\text{starName}} (\text{MovieStar} \times \rho_S(\text{StarsIn})))$$

Translation of correlated select-from-where subqueries

```
SELECT S.movieTitle, M.studioName  
FROM StarsIn S, Movie M  
WHERE S.movieYear >= 2000  
AND S.movieTitle = M.title  
AND EXISTS (SELECT name  
            FROM MovieStar  
            WHERE birthdate=1960 AND name= S.starName)
```

Complete the expression

$$\Pi_{S.movieTitle, M.studioName} \sigma_{S.movieYear >= 2000 \wedge S.movieTitle = M.title} \\ (\rho_M(Movie) \bowtie$$
$$\Pi_{S.movieTitle, S.movieYear, S.starName}$$
$$\sigma_{birthdate=1960 \wedge name=S.starName} (MovieStar \times \rho_S(StarsIn))`$$

Wait !

```
SELECT S.movieTitle, M.studioName  
FROM StarsIn S, Movie M  
WHERE S.movieYear >= 2000  
AND S.movieTitle = M.title  
AND EXISTS (SELECT name  
            FROM MovieStar  
            WHERE birthdate=1960 AND name= S.starName)
```

SQL Result ?

Movie		
=====		
title	studioName	movieYear
DBSA	ULB	2005

StartsIn		
=====		
starName	movieTitle	
Foo	DBSA	

MovieStar		
=====		
name	firstname	birthdate
Foo	Bar	1960
Foo	Baz	1960

Wait !

```
SELECT S.movieTitle, M.studioName  
FROM StarsIn S, Movie M  
WHERE S.movieYear >= 2000  
AND S.movieTitle = M.title  
AND EXISTS (SELECT name  
            FROM MovieStar  
            WHERE birthdate=1960 AND name= S.starName)
```

SQL Result ?

movieTitle	studioName
DBSA	ULB



one result because
subquery only search
for at least one

Movie		
=====		
title	studioName	movieYear
DBSA	ULB	2005

StartsIn	
=====	
starName	movieTitle
Foo	DBSA

MovieStar		
=====		
name	firstname	birthdate
Foo	Bar	1960
Foo	Baz	1960

Wait !

$$\Pi_{S.movieTitle, M.studioName}$$
$$\sigma_{S.movieYear >= 2000 \wedge S.movieTitle = M.title} (\rho_M(Movie) \bowtie$$
$$\Pi_{S.movieTitle, S.movieYear, S.starName} (\sigma_{birthdate = 1960 \wedge name = S.starName} (MovieStar \times \rho_S(StarsIn)))$$

RA Result ?

Movie		
=====		
title	studioName	movieYear

DBSA	ULB	2005
StartsIn		
=====		
starName	movieTitle	

Foo		DBSA
MovieStar		
=====		
name	firstname	birthdate

Foo	Bar	1960
Foo	Baz	1960

Wait !

$$\Pi_{S.movieTitle, M.studioName}$$
$$\sigma_{S.movieYear >= 2000 \wedge S.movieTitle = M.title} (\rho_M(Movie) \bowtie$$
$$\Pi_{S.movieTitle, S.movieYear, S.starName} (\sigma_{birthdate = 1960 \wedge name = S.starName} (MovieStar \times \rho_S(StarsIn)))$$

RA Result ?

movieTitle	studioName
DBSA	ULB
DBSA	ULB

Movie

=====

title studioName movieYear

DBSA ULB 2005

StartsIn

=====

starName movieTitle

Foo DBSA

MovieStar

=====

name firstname birthdate

Foo Bar 1960
Foo Baz 1960

Wait !

$$\Pi_{S.movieTitle, M.studioName}$$
$$\sigma_{S.movieYear >= 2000 \wedge S.movieTitle = M.title}$$
$$\rho_M(Movie) \bowtie$$
$$\Pi_{S.movieTitle, S.movieYear, S.starName}$$
$$\sigma_{birthdate = 1960 \wedge name = S.starName} ($$
$$MovieStar \times \rho_S(StarsIn))$$

RA Result ?

movieTitle studioName

DBSA ULB

DBSA ULB

Only set-based Relations

Movie

=====

title studioName movieYear

DBSA ULB 2005

StartsIn

=====

starName movieTitle

Foo DBSA

MovieStar

=====

name firstname birthdate

Foo Bar 1960

Foo Baz 1960

Flattening Subqueries in **Bag-based** Relations (probably all vendor implementations)

The requirements for flattening into a normal join are:

- There is a uniqueness condition that ensures that the subquery does not introduce any duplicates if it is flattened into the outer query block.
- Each table in the subquery's FROM list (after any view, derived table, or subquery flattening) must be a base table.
- The subquery is not under an OR.
- The subquery is not in the SELECT list of the outer query block.
- The subquery type is EXISTS, IN, or ANY, or it is an expression subquery on the right side of a comparison operator.

Flattening Subqueries in **Bag-based** Relations (probably all vendor implementations)

- There are no aggregates in the SELECT list of the subquery.
- The subquery does not have a GROUP BY clause.
- The subquery does not have an ORDER BY, result offset, or fetch first clause.
- If there is a WHERE clause in the subquery, there is at least one table in the subquery whose columns are in equality predicates with expressions that do not include any column references from the subquery block. These columns must be a superset of the key columns for any unique index on the table. For all other tables in the subquery, the columns in equality predicates with expressions that do not include columns from the same table are a superset of the unique columns for any unique index on the table.

System R: Relational Approach to Database Management

M. M. ASTRAHAN, M. W. BLASGEN, D. D. CHAMBERLIN,
K. P. ESWARAN, J. N. GRAY, P. P. GRIFFITHS,
W. F. KING, R. A. LORIE, P. R. McJONES, J. W. MEHL,
G. R. PUTZOLU, I. L. TRAIGER, B. W. WADE, AND V. WATSON

IBM Research Laboratory

To read before the next lecture. We will discuss it in the lecture. Only read until end of The Optimizer section (unless you fall in love with it)

<https://www.seas.upenn.edu/~zives/cis650/papers/System-R.PDF>

Credits

Many slides are copied from:

- Stijn Vansumeren, Database Systems Architecture course slides.