

The Art of PostgreSQL

Turn Thousands of Lines
of Code into Simple Queries



Postgres Extensibility

Université Libre de Bruxelles



PostgreSQL Extensibility

The Design of POSTGRES


This paper presents the preliminary design of a new database management system, called POSTGRES, that is the successor to the INGRES relational database system. The main design goals of the new system are to:

1. provide better support for complex objects,
2. provide user extendibility for data types, operators and access methods,
3. provide facilities for active databases (i.e., alerters and triggers) and inferencing including forward- and backward-chaining,
4. simplify the DBMS code for crash recovery,
5. produce a design that can take advantage of optical disks, workstations composed of multiple tightly-coupled processors, and custom designed VLSI chips, and
6. make as few changes as possible (preferably none) to the relational model.

The paper describes the query language, programming language interface, system architecture, query processing strategy, and storage system for the new system.

PostgreSQL Extensibility

- *SQL can be Object Oriented
(when using Postgres)*
- *Extensible SQL*
- *Understanding Postgres data
types*
- *And Operators*
- *And Operator Classes*
- *And Postgres Indexing APIs*
- *And Extensions*
- *Then developing a new extension,
in C, re-using Postgres internals
as much as possible*





Extensibility - SQL

Extensibility - SQL

```
select col1, col2 from table where col1 = 'something';
```

for Postgres to return
those, it must first understand
their data types

non decorated
literal, Postgres
will need to find the
data type on its own

Extensibility - SQL

```
SELECT col  
  FROM table  
WHERE stamped > date 'today' - interval '1 day';
```

decoration of
the literal value \Rightarrow decorated
literal

↳ prefixed with
the name of
the data type

Object Oriented
SQL
when using
Postgres

Extensibility - SQL

```
select x,  
       1 + x as "1+",  
       '127.0.0.1'::inet + x as "ip address",  
       date 'today' + x as date  
from (values (0), (1), (2), (3)) as t(x);
```

add an integer
will have a different
effect depending
on the data type
of the other
operand

x	1+	ip address	date
0	1	127.0.0.1	2018-03-22
1	2	127.0.0.2	2018-03-23
2	3	127.0.0.3	2018-03-24
3	4	127.0.0.4	2018-03-25
(4 rows)			

Note: same type of
catalog look up for
a user created
datatype and for
base Postgres type

// different method
depending on the
type of object

Extensibility - SQL

ip4r deals with ranges of IP addresses

```
select iprange, locid  
from geolite.blocks
```

```
where iprange >>= '91.121.37.122';
```

? → new operator for "contains"

*does the range
contain the
IP address*

iprange	locid
91.121.0.0-91.121.159.255	75

(1 row)

Time: 1.220 ms

*comes from
the extension (not native in Postgres)*

Postgres has no hardcoded knowledge, it is sorted in different catalogs that it has to look-up to understand the queries

Operator Classes

→ organize a set of functions as an API allowing for creating indexes

```
select amopopr::regoperator
from pg_opclass c
join pg_am am → access method
on am.oid = c.opcmethod
join pg_amop amop
on amop.amopfamily = c.opcfamily
where opcintype = 'ip4r'::regtype
and am.amname = 'gist';
```

① Here it is looking for the contain operator

amopopr
>=(ip4r, ip4r)
<=(ip4r, ip4r)
>>(ip4r, ip4r)
<<(ip4r, ip4r)
&&(ip4r, ip4r)
=(ip4r, ip4r)

(6 rows)

when reading a query, Postgres does catalog look-ups to recognize and understand the different parts of the query

② look-up to find the code of the operator

③ look-up if there is an index on the column

Extensibility - SQL

```
select id, name, pos,  
       pos <@> point(-6.25,53.34) as miles  
  from pubnames  
order by pos <-> point(-6.25,53.34)  
limit 10;
```

)
"2-nearest neighbors"

Extensibility - Catalogs

```
SELECT CASE WHEN o.oprkind='l' THEN NULL
        ELSE pg_catalog.format_type(o.oprleft, NULL)
        END AS "Left arg type",

CASE WHEN o.oprkind='r' THEN NULL
        ELSE pg_catalog.format_type(o.oprright, NULL)
        END AS "Right arg type",

pg_catalog.format_type(o.oprrresult, NULL) AS "Result type",

oprcode::regprocedure AS function

FROM pg_catalog.pg_operator o
LEFT JOIN pg_catalog.pg_namespace n ON n.oid = o.oprnamespace
LEFT JOIN pg_catalog.pg_type l ON l.oid = o.oprleft

WHERE o.oprname = '+'
      AND l.typname IN ('integer', 'inet', 'date')
      AND pg_catalog.pg_operator_is_visible(o.oid)

ORDER BY 1, 2, 3, 4;
```

Extensibility - Catalogs

Left arg type	Right arg type	Result type	function
date	integer	date	date_pli(date, integer)
date	interval	timestamp without time zone	date_pl_interval(date, interval)
date	time with time zone	timestamp with time zone	datetimetz_pl(date, time with time zone)
date	time without time zone	timestamp without time zone	datetime_pl(date, time without time zone)
inet	bigint	inet	inetpl(inet, bigint)

(5 rows)

Extensibility - Data Types and Indexes

- *Data types*
- *Input/Output Functions*
- *Casts (implicit/explicit)*
- *Operator Classes*
- *Operator Families*

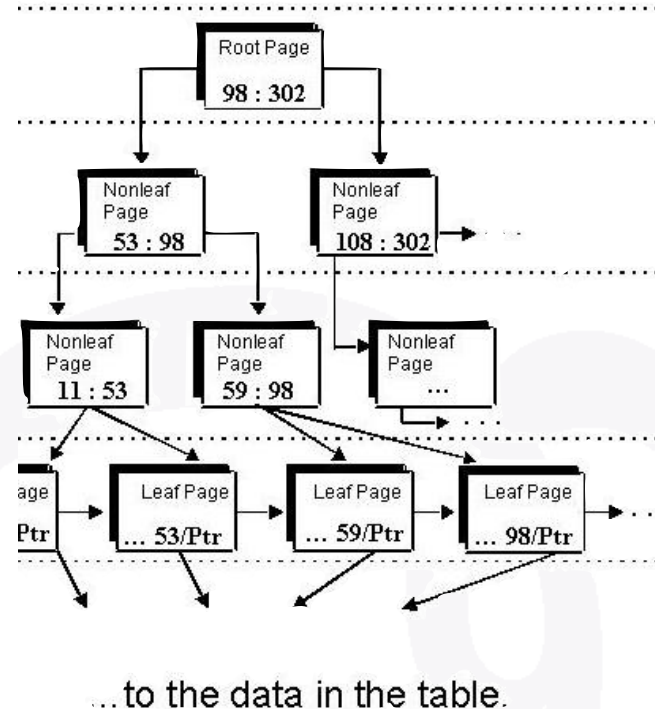




Extensibility - Indexes

Extensibility - Indexes

- *BTree*
- *GiST, Generalized Search Tree*
- *SP-GiST, Space Partitioned Tree*
- *GIN, Generalized Inverted Index*
- *And more (BRIN, bloom, ...)*



Extensibility - Indexes and constraints

```
CREATE TABLE reservation
```

```
(  
    room      text,  
    professor text,  
    during    period,
```

```
EXCLUDE USING gist
```

```
(  
    room with =,  
    during with &&  
)
```

```
);
```

```
CREATE TABLE circles (  
    c circle,  
    EXCLUDE USING gist (c WITH &&)  
);
```

overlapping

*) twice the same room
with overlapping periods*

Indexes - BTree

- *Built for Speed*
- *Unique concurrency tricks*
- *Balanced*
- *Support Function: cmp*
- *Operators: <= < = > >=*

total ordering

// strcmp

1 | 3 |

2 |

1 | 3 | 4 |

2 | 4 |

1 | 3 | 5 |

2 | 4 |

1 | 3 | 5 | 6 |

generalized
⇒ slower than B-tree
↑ that are focused


Indexes - GiST

can't be implemented
using B-tree
because not
strictly greater, ...

- Built for comfort
- Balanced
- Support Functions:
 - consistent, same, union
 - penalty, pick split
 - compress, decompress
- Operators:

@> <@ && = &< &> <<| ...

R-tree Hierarchy



Indexes - GIN

- *Built for Text Search, arrays, JSON*

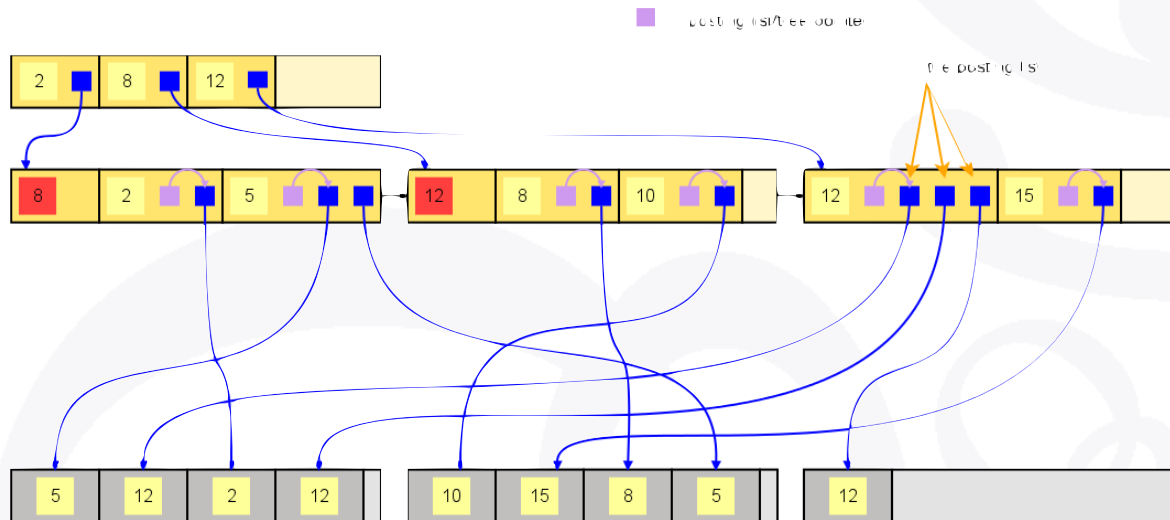
- *Balanced*

- *Support Functions:*

- compare, consistent
- extractValue, extractQuery

- *Operators:*

@> <@ && =






Postgres Extensions

Geolocation with Postgres

- *Using ip4r extension*
- *GiST Indexes*
- *kNN searches*



Ip address ranges



Ip address ranges

```
table geolite.blocks limit 10;
```

iprange		locid
-----+-----		
1.0.0.0/24		17
1.0.1.0-1.0.3.255		49
1.0.4.0/23		14409
1.0.6.0/23		17
1.0.8.0/21		49
1.0.16.0/20		14614
1.0.32.0/19		47667
1.0.64.0/18		111
1.0.128.0-1.0.147.255		209
1.0.148.0/24		22537

(10 rows)

Geolocation in Postgres

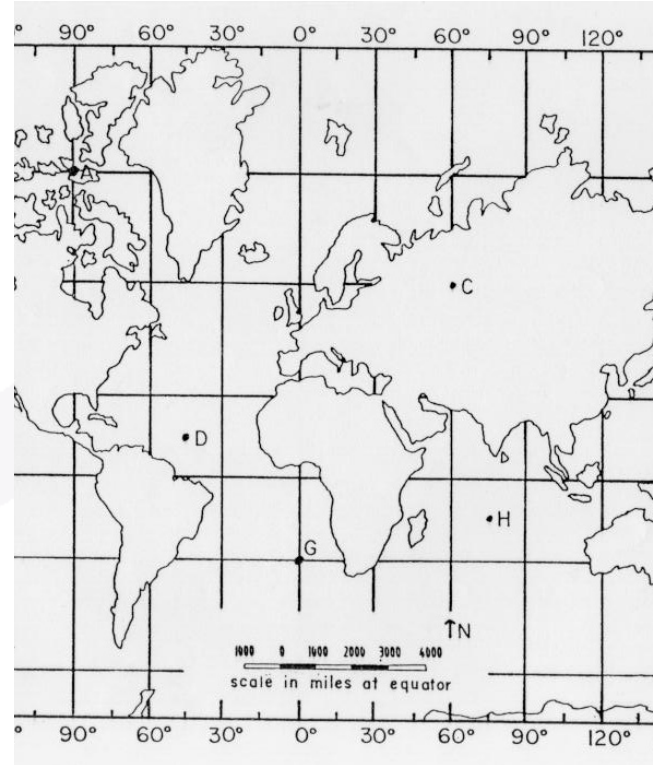
```
select *  
  from geolite.blocks  
 join geolite.location  
    using(loid)  
where iprange  
      >= '74.125.195.147';
```

```
-[ RECORD 1 ]-----  
loid          | 2703  
iprange       | 74.125.189.24-74.125.255.255  
country       | US  
region        | CA  
city          | Mountain View  
postalcode    | 94043  
location      | (-122.0574,37.4192)  
metrocode     | 807  
areacode      | 650
```

Time: 1.335 ms

Geolocation - type point

```
CREATE TABLE pubnames  
(  
    id    bigint,  
    pos   POINT,  
    name  text  
);
```



How far is the nearest pub

```
create extension cube;  
create extension earthdistance;
```

```
select name,  
       pos <@> point(-6.25,53.34) miles  
  from pubnames  
limit 3;
```

name		pos
-----	+	-----
Ned's		(-6.25,53.34)
Sub Lo		(-6.25,53.34)
O'Neil		(-6.25,53.34)
(3 rows)		

Time: 0.849 ms

Geolocation: ip4r meets earthdistance



Ten nearest pubs

```
with geoloc as
(
  select location as l
    from location
   join blocks using(loid)
  where iprange
        >>=
        '212.58.251.195'
)
select name,
       pos <@> l miles
  from pubnames, geoloc
order by pos <-> l limit 10;
```

name	miles
Blue Anchor	0.299
Dukes Head	0.360
Blue Ball	0.337
Bell (aka The Rat)	0.481
on the Green	0.602
Fox & Hounds	0.549
Chequers	0.712
Sportsman	1.377
Kingswood Arms	1.205
Tattenham Corner	2.007
(10 rows)	

Time: 3.275 ms




Create Extension

Extension base36

<https://github.com/dimitri/base36>

New integer data type: base36

- *Internally a bigint*
- *Visually, a base36 number*
- *Re-use Postgres internals*
- *Provide new I/O functions*



A new integer data type: base36

`create extension base36;`

i		x
-+-----		
0		0
1		1
2		2
3		3
4		4
5		5
6		6
7		7
8		8
9		9
10		A

i		x
-----+-----		
10000		7PS
10001		7PT
10002		7PU
10003		7PV
10004		7PW
10005		7PX
10006		7PY
10007		7PZ
10008		7Q0
10009		7Q1
10010		7Q2

i		x
-----+-----		
100000000		1NJCHS
100000001		1NJCHT
100000002		1NJCHU
100000003		1NJCHV
100000004		1NJCHW
100000005		1NJCHX
100000006		1NJCHY
100000007		1NJCHZ
100000008		1NJCI0
100000009		1NJCI1
100000010		1NJCI2

Using the base36 extension

```
create extension base36;

create table demo(i bigint, x base36);

insert into demo(i, x)
  select n, n::bigint
  from generate_series(0, 10) t(n);

insert into demo(i, x)
  select n, n::bigint
  from generate_series(10000, 10010) t(n);

insert into demo(i, x)
  select n, n::bigint
  from generate_series(1000000000, 1000000010) t(n);

create index on demo(x);
```

Extensibility - I/O functions

```
CREATE OR REPLACE FUNCTION base36_in(cstring) RETURNS base36  
AS '$libdir/base36'  
LANGUAGE C IMMUTABLE STRICT;
```

```
CREATE OR REPLACE FUNCTION base36_out(base36) RETURNS cstring  
AS '$libdir/base36'  
LANGUAGE C IMMUTABLE STRICT;
```

Extensibility - I/O functions

```
CREATE OR REPLACE FUNCTION base36_recv(internal) RETURNS base36  
AS '$libdir/base36'  
LANGUAGE C IMMUTABLE STRICT;
```

```
CREATE OR REPLACE FUNCTION base36_send(base36) RETURNS bytea  
AS '$libdir/base36'  
LANGUAGE C IMMUTABLE STRICT;
```

Extensibility - I/O functions

```
#include "postgres.h"
#ifndef PG_VERSION_NUM
#error "Unsupported too old PostgreSQL version"
#endif

#if PG_VERSION_NUM / 100 != 903 \
    && PG_VERSION_NUM / 100 != 904
#error "Unknown or unsupported PostgreSQL version"
#endif

PG_MODULE_MAGIC;
```

Extensibility - I/O functions

```
static inline
base36 base36_from_str(const char *str)
{
    /* ... C code here ... */
}

static inline
char *base36_to_str(base36 c)
{
    /* ... C code here ... */
}
```

Extensibility - I/O functions

```
Datum base36_in(PG_FUNCTION_ARGS);  
Datum base36_out(PG_FUNCTION_ARGS);  
Datum base36_recv(PG_FUNCTION_ARGS);  
Datum base36_send(PG_FUNCTION_ARGS);  
Datum base36_cast_to_text(PG_FUNCTION_ARGS);  
Datum base36_cast_from_text(PG_FUNCTION_ARGS);  
Datum base36_cast_to_bigint(PG_FUNCTION_ARGS);  
Datum base36_cast_from_bigint(PG_FUNCTION_ARGS);
```


Extensibility - I/O functions

```
PG_FUNCTION_INFO_V1(base36_in);
```

```
Datum base36_in(PG_FUNCTION_ARGS)
{
    char *str = PG_GETARG_CSTRING(0);
    PG_RETURN_INT64(base36_from_str(str));
}
```

```
PG_FUNCTION_INFO_V1(base36_out);
```

```
Datum base36_out(PG_FUNCTION_ARGS)
{
    base36 c = PG_GETARG_INT64(0);
    PG_RETURN_CSTRING(base36_to_str(c));
}
```

Extensibility - I/O functions

```
CREATE TYPE base36
```

```
(
```

```
    INPUT      = base36_in,  
    OUTPUT     = base36_out,  
    RECEIVE    = base36_recv,  
    SEND       = base36_send,  
    LIKE       = bigint,  
    CATEGORY   = 'N'
```

```
);
```

```
COMMENT ON TYPE base36
```

```
    IS 'bigint written in base36: [0-9A-Z]+';
```

Extensibility - I/O functions

```
CREATE FUNCTION text(base36)
RETURNS text
AS '$libdir/base36',
'base36_cast_to_text'
LANGUAGE C IMMUTABLE STRICT;
```

```
CREATE CAST (text as base36)
WITH FUNCTION base36(text)
AS IMPLICIT;
```

```
CREATE CAST (base36 as text)
WITH FUNCTION text(base36);
```

```
CREATE CAST (bigint as base36)
WITHOUT FUNCTION
AS IMPLICIT;
```

```
CREATE CAST (base36 as bigint)
WITHOUT FUNCTION
AS IMPLICIT;
```

Reuse Postgres Internals

```
CREATE OR REPLACE FUNCTION base36_eq(base36, base36)  
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8eq';
```

```
CREATE OR REPLACE FUNCTION base36_ne(base36, base36)  
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8ne';
```

```
CREATE OR REPLACE FUNCTION base36_lt(base36, base36)  
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8lt';
```

```
CREATE OR REPLACE FUNCTION base36_le(base36, base36)  
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8le';
```

Reuse Postgres Internals

```
CREATE OR REPLACE FUNCTION base36_gt(base36, base36)  
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8gt';
```

```
CREATE OR REPLACE FUNCTION base36_ge(base36, base36)  
RETURNS boolean LANGUAGE internal IMMUTABLE AS 'int8ge';
```

```
CREATE OR REPLACE FUNCTION base36_cmp(base36, base36)  
RETURNS integer LANGUAGE internal IMMUTABLE AS 'btint8cmp';
```

Extensibility - Operators

```
CREATE OPERATOR = (  
    LEFTARG = base36,  
    RIGHTARG = base36,  
    PROCEDURE = base36_eq,  
    COMMUTATOR = '=',  
    NEGATOR = '<>',  
    RESTRICT = eqsel,  
    JOIN = eqjoinsel  
);  
  
COMMENT ON OPERATOR =(base36, base36) IS 'equals?';
```

Extensibility - Operator Class

```
CREATE OPERATOR CLASS btree_base36_ops  
DEFAULT FOR TYPE base36 USING btree  
AS
```

OPERATOR	1 < ,
OPERATOR	2 <= ,
OPERATOR	3 = ,
OPERATOR	4 >= ,
OPERATOR	5 > ,
FUNCTION	1 base36_cmp(base36, base36);